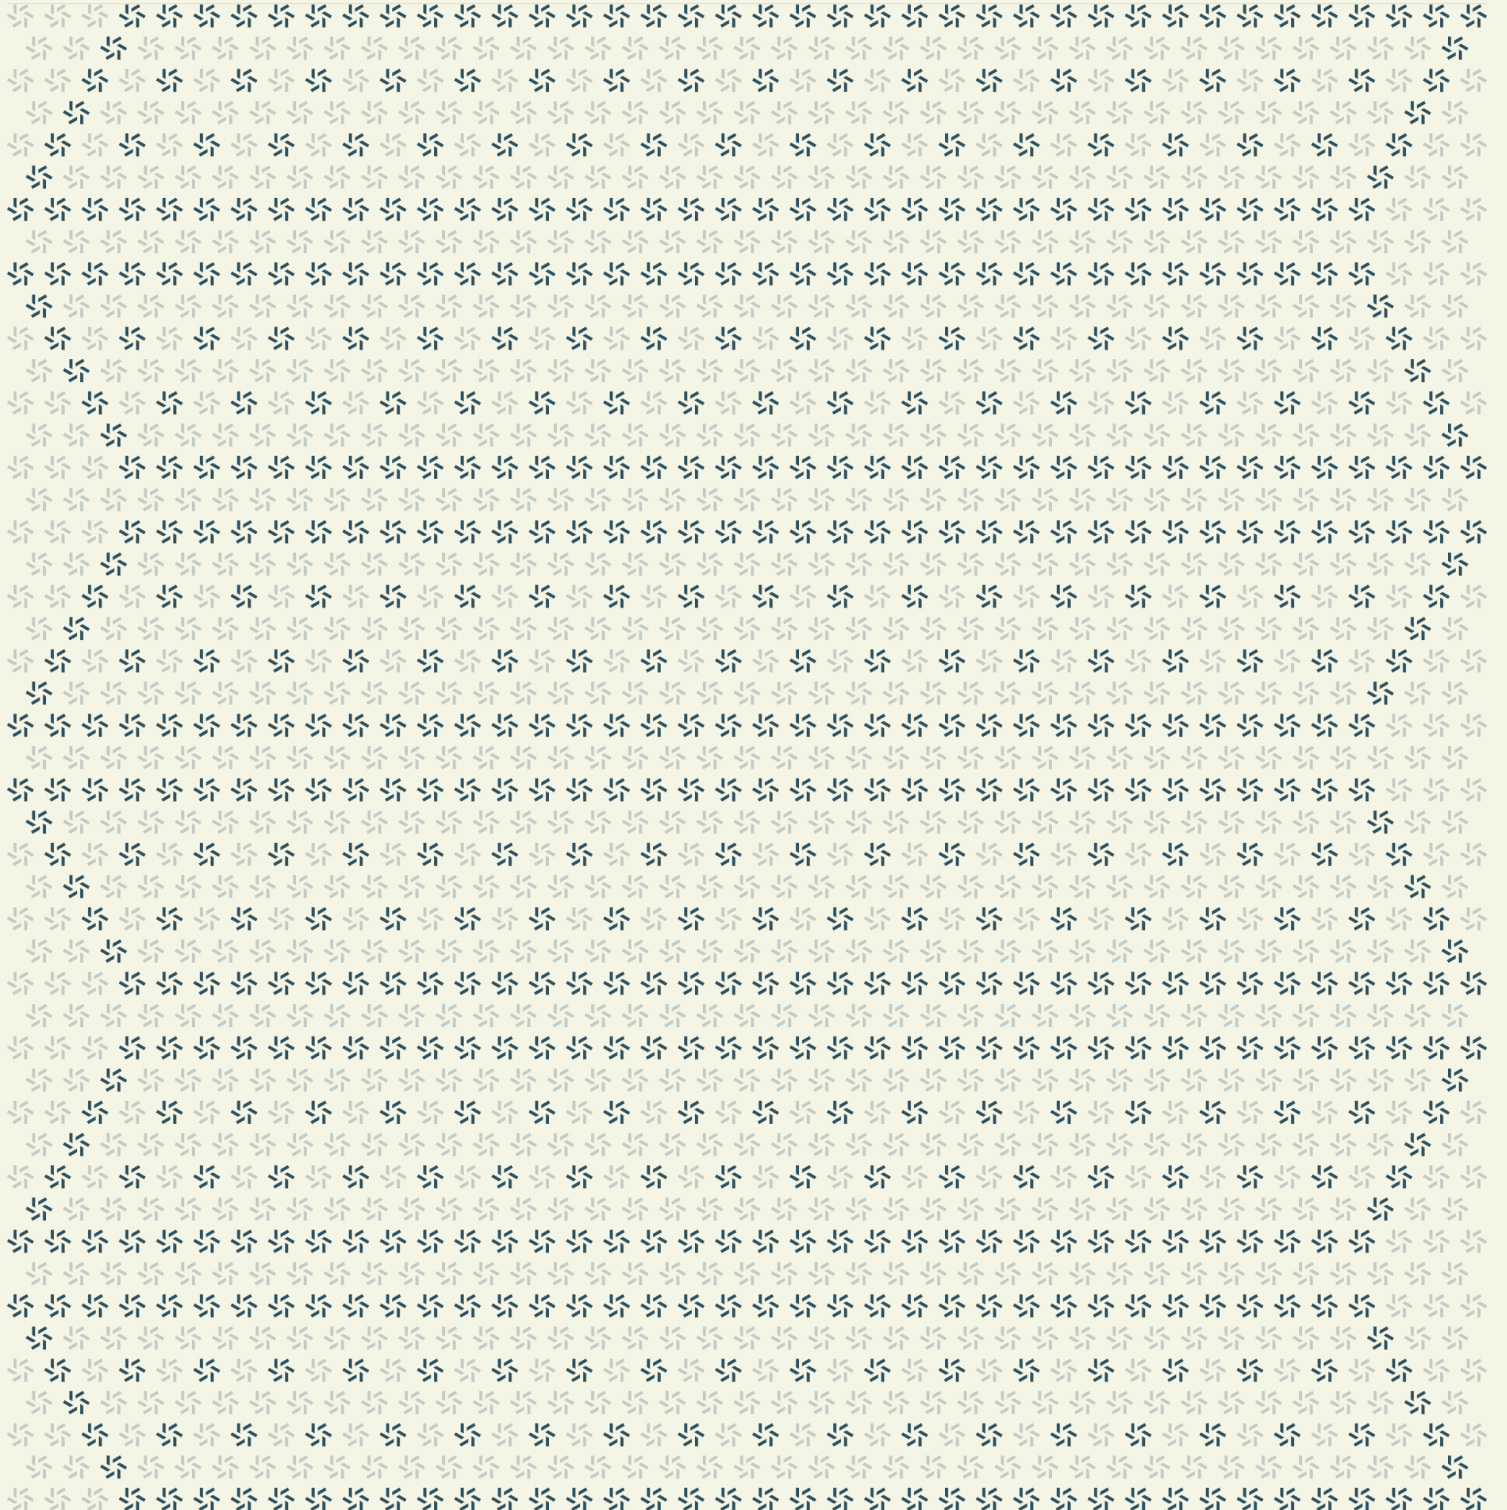


January 17, 2025

Pyth Lazer

Solana Application Security Assessment



Contents

About Zellic	3
<hr/>	
1. Overview	3
1.1. Executive Summary	4
1.2. Goals of the Assessment	4
1.3. Non-goals and Limitations	4
1.4. Results	4
<hr/>	
2. Introduction	5
2.1. About Pyth Lazer	6
2.2. Methodology	6
2.3. Scope	8
2.4. Project Overview	8
2.5. Project Timeline	9
<hr/>	
3. Detailed Findings	9
3.1. Signature bypass	10
3.2. Migration applicable to already migrated storage account	13
<hr/>	
4. Threat Model	14
4.1. Program: Lazer	15
<hr/>	
5. Assessment Results	19
5.1. Disclaimer	20

About Zellic

Zellic is a vulnerability research firm with deep expertise in blockchain security. We specialize in EVM, Move (Aptos and Sui), and Solana as well as Cairo, NEAR, and Cosmos. We review L1s and L2s, cross-chain protocols, wallets and applied cryptography, zero-knowledge circuits, web applications, and more.

Prior to Zellic, we founded the [#1 CTF \(competitive hacking\) team](#) worldwide in 2020, 2021, and 2023. Our engineers bring a rich set of skills and backgrounds, including cryptography, web security, mobile security, low-level exploitation, and finance. Our background in traditional information security and competitive hacking has enabled us to consistently discover hidden vulnerabilities and develop novel security research, earning us the reputation as the go-to security firm for teams whose rate of innovation outpaces the existing security landscape.

For more on Zellic's ongoing security research initiatives, check out our website zellic.io and follow [@zellic_io](#) on Twitter. If you are interested in partnering with Zellic, contact us at hello@zellic.io.



1. Overview

1.1. Executive Summary

Zellic conducted a security assessment for Pyth Data Association from December 11th to December 12th, 2024. During this engagement, Zellic reviewed Pyth Lazer's code for security vulnerabilities, design issues, and general weaknesses in security posture.

1.2. Goals of the Assessment

In a security assessment, goals are framed in terms of questions that we wish to answer. These questions are agreed upon through close communication between Zellic and the client. In this assessment, we sought to answer the following questions:

- Is the price feed signature verification functionality implemented correctly and securely?
 - Is the administrative functionality implemented correctly and securely?
-

1.3. Non-goals and Limitations

We did not assess the following areas that were outside the scope of this engagement:


- Front-end components
 - Back-end components
 - Infrastructure relating to the project
 - Key custody
-

1.4. Results

During our assessment on the scoped Pyth Lazer programs, we discovered two findings. One critical issue was found. One was of low impact.

Breakdown of Finding Impacts

Impact Level	Count
<div>Critical</div>	1
<div>High</div>	0
<div>Medium</div>	0
<div>Low</div>	1
<div>Informational</div>	0



2. Introduction

2.1. About Pyth Lazer

Pyth Data Association contributed the following description of Pyth Lazer:

Lazer is a blazing-fast oracle with focus on latency, and efficiency. It streams price updates in a flexible format for on-chain consumption.

2.2. Methodology

During a security assessment, Zellic works through standard phases of security auditing, including both automated testing and manual review. These processes can vary significantly per engagement, but the majority of the time is spent on a thorough manual review of the entire scope.

Alongside a variety of tools and analyzers used on an as-needed basis, Zellic focuses primarily on the following classes of security and reliability issues:

Basic coding mistakes. Many critical vulnerabilities in the past have been caused by simple, surface-level mistakes that could have easily been caught ahead of time by code review. Depending on the engagement, we may also employ sophisticated analyzers such as model checkers, theorem provers, fuzzers, and so on as necessary. We also perform a cursory review of the code to familiarize ourselves with the programs.

Business logic errors. Business logic is the heart of any smart contract application. We examine the specifications and designs for inconsistencies, flaws, and weaknesses that create opportunities for abuse. For example, these include problems like unrealistic tokenomics or dangerous arbitrage opportunities. To the best of our abilities, time permitting, we also review the contract logic to ensure that the code implements the expected functionality as specified in the platform's design documents.

Integration risks. Several well-known exploits have not been the result of any bug within the contract itself; rather, they are an unintended consequence of the contract's interaction with the broader DeFi ecosystem. Time permitting, we review external interactions and summarize the associated risks: for example, flash loan attacks, oracle price manipulation, MEV/sandwich attacks, and so on.

Code maturity. We look for potential improvements in the codebase in general. We look for violations of industry best practices and guidelines and code quality standards. We also provide suggestions for possible optimizations, such as gas optimization, upgradability weaknesses, centralization risks, and so on.

For each finding, Zellic assigns it an impact rating based on its severity and likelihood. There is no hard-and-fast formula for calculating a finding's impact. Instead, we assign it on a case-by-case basis based on our judgment and experience. Both the severity and likelihood of an issue affect its impact. For instance, a highly severe issue's impact may be attenuated by a low likelihood.

We assign the following impact ratings (ordered by importance): Critical, High, Medium, Low, and Informational.

Zellic organizes its reports such that the most important findings come first in the document, rather than being strictly ordered on impact alone. Thus, we may sometimes emphasize an "Informational" finding higher than a "Low" finding. The key distinction is that although certain findings may have the same impact rating, their *importance* may differ. This varies based on various soft factors, like our clients' threat models, their business needs, and so on. We aim to provide useful and actionable advice to our partners considering their long-term goals, rather than a simple list of security issues at present.

2.3. Scope

The engagement involved a review of the following targets:

Pyth Lazer Programs

Type	Rust
Platform	Solana
Target	Pyth monorepo
Repository	https://github.com/pyth-network/pyth-crosschain ↗
Version	2faddf96bc22e1aedc1f72f929cb1a64b520db6a
Programs	lib.rs signature.rs

2.4. Project Overview

Zellic was contracted to perform a security assessment for a total of two person-days. The assessment was conducted by two consultants over the course of two calendar days.

Contact Information

The following project managers were associated with the engagement:

Jacob Goreski
↗ Engagement Manager
jacob@zellic.io ↗

Chad McDonald
↗ Engagement Manager
chad@zellic.io ↗

The following consultants were engaged to conduct the assessment:

Jasraj Bedi
↗ Co-founder
jazzy@zellic.io ↗

Filippo Cremonese
↗ Engineer
fcremo@zellic.io ↗

2.5. Project Timeline

The key dates of the engagement are detailed below.

December 11, 2024 Start of primary review period

December 12, 2024 End of primary review period

3. Detailed Findings

3.1. Signature bypass

Target	Lazer		
Category	Coding Mistakes	Severity	Critical
Likelihood	Low	Impact	Critical

Description

verify_message incorrectly assumes that the layout of the calldata passed to the VerifyMessage instruction is as follows:

Offset	Size	Handler argument name/Significance
0	8	ANCHOR HEADER
8	4	message_data length (4+64+32+2+psz)
12		message_data
12	4	SOLANA_FORMAT_MAGIC
16	64	signature
80	32	pubkey
112	2	payload size (psz)
114	psz	payload
114+psz	2	ed25519_instruction_index
116+psz	1	signature_index
117+psz	2	message_offset

Under this assumption, message_offset has a value of 12, referencing the offset of the data of the message_data vector in the raw encoding of the calldata for the VerifyMessage instruction.

The verify_message function performs several checks on the data contained in message_data using offsets computed assuming the storage layout outlined above.

However, if the calldata layout and message_offset are adjusted in the following way:

Offset	Size	Handler argument name/Significance
0	8	ANCHOR HEADER
8	4	message_data length (4+64+32+2+psz)
12		message_data
12	4	SOLANA_FORMAT_MAGIC
16	64	<UNUSED>
80	32	PK1
112	2	payload1 size (psz1)
114	psz1	payload1
114+psz1	4	<UNUSED>

118+psz1	64	valid signature for PK2
182+psz1	32	PK2
214+psz1	2	payload2 size (psz2)
216+psz1	psz	payload2
216+psz1+psz2	2	ed25519_instruction_index
218+psz1+psz2	1	signature_index
219+psz1+psz2	2	message_offset (114+psz1)

The `verify_message` function will read the public key, payload size and payload appearing earlier in the calldata, while the `ec25519` instruction will reference the latter instances.

Thus, `verify_message` will incorrectly check that the PK1 public key appears in the list of trusted signers, while the `ec25519` instruction will require a signature from a different PK2 public key. Additionally, the signed payload and the payload returned by `verify_message` also differ arbitrarily both in their content and their size.

Impact

If exploited, this issue allows to completely bypass signature verification of a price feed update, allowing to forge a price feed update with an arbitrary payload and that appears to originate from any address in the set of trusted signers.

The exploitability of this issue is conditional on the program invoking the `VerifyMessage` instruction. The attacker needs to have control over the `message_offset` provided to `VerifyMessage`. The official integration examples provided by Pyth Data Association did not allow the end user to specify the offset provided to the `VerifyMessage` instruction. These preconditions lower the likelihood of the vulnerability. However, Lazer integrations are not required to use the examples provided by Pyth Data Association, and there seems to be some more advanced and reasonable scenarios where a program integrating with Lazer could give an attacker control over the offset argument.

Recommendations

It appears to be possible to entirely remove the `message_offset` argument, fixing its value to 12, which would constrain the data referenced by the `ec25519` program to be aligned with the start of the `message_data` argument of the `VerifyMessage` instruction. Note that the offset used by the `ec25519` program would still need to be checked to be 12 if removing the `message_offset` argument.

Remediation

This issue has been acknowledged by Pyth Data Association, and a fix was implemented in [PR #2250](#).

The PR added a check requiring the `message_data` processed by the `verify_message` function to match the data considered by the `ec25519` program.

Additionally, the `message_offset` argument indicating the start of the message data processed by the `ec25519` program was removed.

The code could not be simplified further by requiring a stricter layout for the data processed by the `ec25519` program, as Pyth Data Association indicated this flexibility is a requirement.

3.2. Migration applicable to already migrated storage account

Target	Lazer		
Category	Coding Mistakes	Severity	Low
Likelihood	Medium	Impact	Low

Description

`migrate_from_0_1_0` allows the `top_authority` to migrate the storage account from the older `StorageV010` model to the new `Storage` data model. This resizes the account, migrates the `trusted_signers`, and sets the `treasury` based on a value passed in by the `top_authority`. The migration requires a few conditions to be true, but one of the most important bits is that the anchor signature is equivalent to the `Storage` account's signature. Notably, this means that the migration does *not* change the anchor signature, because the name of the type for the account has not changed.

As the preliminary checks that allow the migration to go forward are true both before and after the migration, nothing prevents the `top_authority` from rerunning the migration again.

Impact

The `top_authority` could reset the `treasury` and/or putting the `trusted_signers` in a potentially corrupted state. This could be done by causing the `num_trusted_signers` to be set to an invalid value, breaking the signature validation logic.

Recommendations

As Anchor's discriminant is primarily dependant on the name of the type of the account, we recommend changing the name of the account from `Storage` to a distinct name such as `StorageV2`. The migration step would validate that the account had a discriminant expected for `Storage`, but would eventually write `StorageV2`'s discriminant to the account. This would prevent subsequent migrations as the discriminant of the account would change.

Remediation

This issue has been acknowledged by Pyth Data Association, and a fix was implemented in [PR #2250](#).

The PR added a check requiring the size of the `Storage` account to match the size of the legacy layout. Since the legacy and new layouts differ in size, this prevents the migration from being performed more than once.

Pyth Data Association subsequently removed the migration feature entirely in commit [800bb1aa](#).

4. Threat Model

This provides a full threat model description for various functions. As time permitted, we analyzed each function in the programs and created a written threat model for some critical functions. A threat model documents a given function's externally controllable inputs and how an attacker could leverage each input to cause harm.

Not all functions in the audit scope may have been modeled. The absence of a threat model in this section does not necessarily suggest that a function is safe.

4.1. Program: Lazer

initialize

This instruction can be used to create and initialize the singleton storage account. The instruction handler function does not enforce access control, and therefore it could technically be invoked by any caller. However, the storage account being created has a fixed seed. For this reason, `initialize` can only be called once. While this does not directly prevent someone to race to call the instruction and initialize the storage account, this action would be detected and could be addressed by simply redeploying and reinitializing a new instance of the program.

The storage account `top_authority` and `treasury` fields are initialized to public key values given as arguments to the `initialize` instruction, and `single_update_fee_in_lamports` is initialized to 1.

Arguments

- `top_authority`: public key value used to initialize the `top_authority` field in the newly created storage account
- `treasury`: public key value used to initialize the `treasury` field in the newly created storage account

Accounts

- `payer`: account paying for the fees needed to create the storage account
 - Anchor flags: `signer, mut`
- `storage`: new account being initialized
 - Anchor flags: `init`
 - Payer: `payer`
 - Seed: `STORAGE_SEED` (fixed constant)
 - Space: manually specified

Tests

Happy path behavior is implicitly tested by the testsuite setup routines.

`migrate_from_0_1_0`

This instruction can be used to migrate the singleton storage account in-place from the legacy layout to the current storage layout.

The instruction handler takes the storage account as a raw `AccountInfo`, manually verifies the discriminant and deserializes it according to the legacy layout.

It then ensures that the `top_authority` signer account matches the `top_authority` field of the deserialized storage account.

The storage account is reallocated, increasing its size to account for the additional space needed by the newer Storage account type.

Finally, the function reinitializes the account, by instantiating an instance of the new Storage account type and serializing it manually into the storage account.

This function does not prevent to migrate an already migrated storage account, allowing to corrupt its data. Refer to finding [3.2](#) ↗ for a more in depth discussion.

Arguments

- `treasury`: public key value used to initialize the treasury field of the migrated storage account

Accounts

- `top_authority`: Signer account giving authorizing the migration
 - Anchor flags: signer
 - Checks: must match the `top_authority` field of the storage account when decoded using the legacy layout
- `storage`: storage account to be migrated; taken as a raw `AccountInfo`
 - Anchor flags: mut
 - Seed: `STORAGE_SEED` (fixed constant)

Tests

Positive cases

- ☑ A legacy Storage account can be migrated successfully, with the new Storage account layout correctly reflecting the previous configuration

Negative cases

- ☒ Rejects attempts to migrate the Storage account multiple times
 - Note: this test was added as part of the remediations
- ☐ Rejects attempts to migrate the Storage account by an unauthorized account

update

This instruction allows to manage the set of trusted signers. It can only be invoked with a signature from the `top_authority` account.

Arguments

- `trusted_signer`: Public key value of the trusted signer to be added/removed/modified
- `expires_at`: Expiration date to be associated to the trusted signer; if zero, signifies the trusted signer should be removed from the signer set

Accounts

- `top_authority`: Signer account giving authorizing the migration
 - Anchor flags: `signer`
 - Checks: must match the `top_authority` field of the storage account
- `storage`: storage account to be migrated; taken as a raw `AccountInfo`
 - Anchor flags: `mut`
 - Seed: `STORAGE_SEED` (fixed constant)
 - Notes: the storage account must have been migrated, deserialization would fail for a legacy-layout account due to mismatching size

Tests

Positive cases

- ☒ Allows the authorized administrator address to add a trusted signer
 - Note: this is implicitly tested by the test setup procedures
- ☐ Allows the authorized administrator address to remove a trusted signer
- ☐ Allows the authorized administrator address to update a trusted signer

Negative cases

- ☐ Rejects calls from unauthorized signers

verify_message

This instruction can be used to verify the legitimacy of a price feed update.

The instruction handler collects the required fee amount specified by the `single_update_fee_in_lamports` of the storage account, transferring it from the payer account to the treasury account.

It then verifies that the message data was correctly signed by ensuring that the transaction contains an instruction invoking the `ed25519` system program which ensures a price feed update was validly signed by one of the currently active trusted signers.

Verification of the `ed25519` instruction requires several checks, including:

- ensuring that the instruction indices used to retrieve the signature, public key, and signed message by the `ed25519` instruction refer to the current `verify_message` instruction
- ensuring that the `ed25519` instruction index is lower than the current `verify_message` instruction
- (implicitly) ensuring that at least one signature is present
- ensuring that the signed data starts with the magic value `SOLANA_FORMAT_MAGIC_LE` (2182742457) to disambiguate messages that may have been signed for other purposes
- ensuring the signer pubkey appears in the set of trusted signers
- ensuring the signer trust has not expired
- ensuring that the offsets for the signature, signer pubkey, and payload provided to the `ed25519` instruction are consistent with the offsets used by `verify_message`
- ensuring that the payload length is consistent

Note that the offset checks are implemented incorrectly, allowing a complete bypass of the signature verification checks. Refer to finding [3.1](#) for a more in-depth discussion of the issue.

Arguments

- `message_data`: signed message that is being verified
- `ed25519_instruction_index`: index of the `ed25519_program` instruction within the transaction. This instruction must precede the current instruction
- `signature_index`: index of the signature within the inputs to the `ed25519_program`
- `message_offset`: offset of the signed message within the input data for the current instruction

Accounts

- payer: account paying for the price feed verification fees
 - Anchor flags: signer, mut
- storage: account storing Lazer configuration
 - Seed: `STORAGE_SEED` (fixed constant)

- Notes: the storage account must have been migrated, deserialization would fail for a legacy-layout account due to mismatching size
- treasury: Treasury token account used to collect fees
 - Checks: Must match the treasury field of the storage account
- instructions_sysvar: Sysvar account used for instruction introspection
 - Notes: this is taken as a raw AccountInfo; the account ID is checked by the Solana SDK functions that receive it as an argument

Tests

Positive cases

- ☒ Allows to verify a payload signed by an authorized signer

Negative cases

- ☐ Rejects payloads with an invalid signature (invocation of a program different than ec25519)
- ☐ Rejects payloads with an invalid signature (not originating from a trusted signer)
- ☐ Rejects payloads with an invalid signature (originating from an expired trusted signer)
- ☐ Rejects invalid payloads (not starting with the correct magic)
- ☒ Rejects all combinations of invalid offsets
- Note: this test was added as part of the remediations

5. Assessment Results

At the time of our assessment, the reviewed code was not deployed to the Solana Mainnet.

During our assessment on the scoped Pyth Lazer programs, we discovered two findings. One critical issue was found. One was of low impact.

5.1. Disclaimer

This assessment does not provide any warranties about finding all possible issues within its scope; in other words, the evaluation results do not guarantee the absence of any subsequent issues. Zellic, of course, also cannot make guarantees about any code added to the project after the version reviewed during our assessment. Furthermore, because a single assessment can never be considered comprehensive, we always recommend multiple independent assessments paired with a bug bounty program.

For each finding, Zellic provides a recommended solution. All code samples in these recommendations are intended to convey how an issue may be resolved (i.e., the idea), but they may not be tested or functional code. These recommendations are not exhaustive, and we encourage our partners to consider them as a starting point for further discussion. We are happy to provide additional guidance and advice as needed.

Finally, the contents of this assessment report are for informational purposes only; do not construe any information in this report as legal, tax, investment, or financial advice. Nothing contained in this report constitutes a solicitation or endorsement of a project by Zellic.